

Advanced Concepts in Machine Learning for Energy

Time: March 26, 27, 30, April 1, 3

1:00 pm - 5:30 pm (in-person students)

(10 more days for distance students)

Location: GERB 315 and Zoom

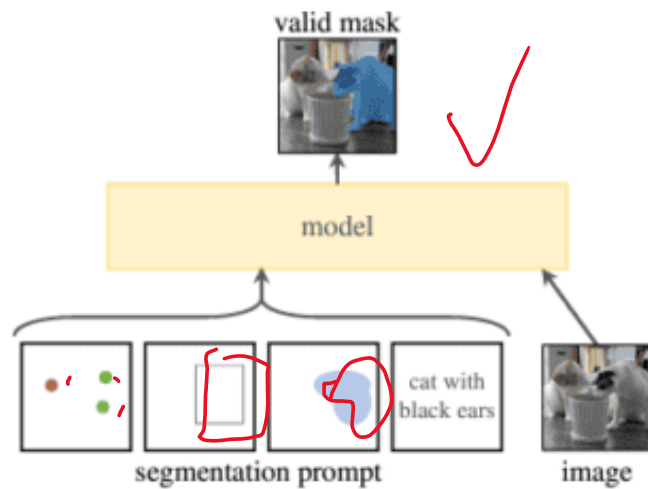
Lecture	Topic	Quiz
1	Neural network fundamentals and continuous optimization	
2	Image classification, object detection, and pixel segmentation	#1
3	Autoencoder, autodecoder, and generative adversarial network (GAN)	#2
4	Attention mechanism, transformers (GPT)	#3
5	Final project Q&A, advanced topics review	#4

Lecture 3


- Quiz #2 on Canvas (1:00-1:30 pm)
- Advanced computer vision (1:30-2:00 pm)
- Lab 3a (2:00-2:30 pm)
- Generative AI (2:30-3:10 pm)
- Lab 3b(3:10-3:40 pm)
- Anomaly Detection using Autoencoders (3:40-4:00 pm)
- **Lab 3c** (4:00-4:30 pm)
- GAN(4:30-5:30 pm)
- **Lab 3d** (5:00-5:30 pm)

Segment anything model (SAM)

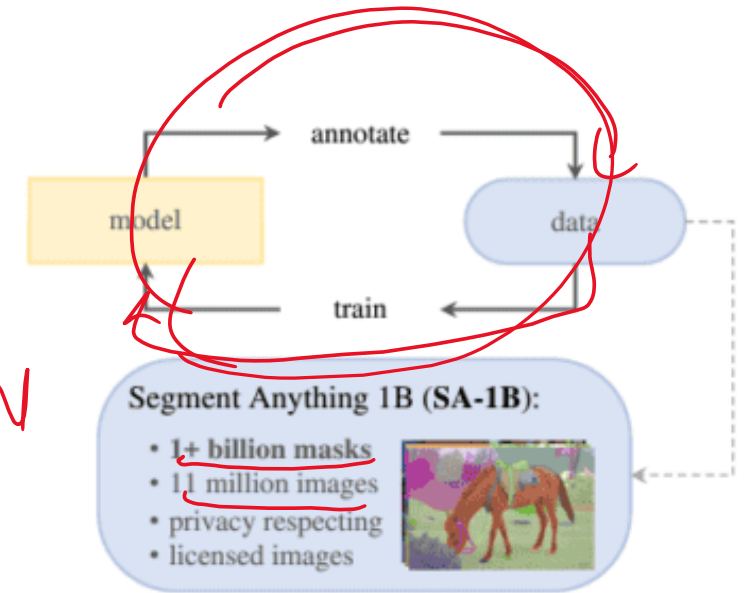
- Foundational model for segmentation
- <https://aidemos.meta.com/segment-anything/>



(a) **Task:** promptable segmentation



(b) **Model:** Segment Anything Model (SAM)




(c) **Data:** data engine (top) & dataset (bottom)

SAM 3D body




Video processing




The first-ever stop-motion film was created by Eadweard Muybridge in 1878. Called *The Horse in Motion*, it's an 11-frame clip that was created to settle a debate over whether all four hooves of a horse were off the ground at the same time when it gallops.

3D reconstruction

- COLMAP
- Scale-invariant feature transform (SIFT)




Feature matching



7-8 pairs

$\begin{bmatrix} u_1 \\ v_1 \end{bmatrix}$

$\begin{bmatrix} u_2 \\ v_2 \end{bmatrix}$



$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} a_u & 0 & u_o \\ 0 & a_v & v_o \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} r_1 & r_2 & r_3 \\ r_4 & r_5 & r_6 \\ r_7 & r_8 & r_9 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} t_1 \\ t_2 \\ t_3 \\ 1 \end{bmatrix} \begin{bmatrix} x_w \\ y_w \\ z_w \\ 1 \end{bmatrix}$$

$$P = K[R|T]$$






Figure 3.6: *Incremental structure from motion.* My incremental structure from motion approach reconstructs the scene a few cameras at a time. This sequence of images shows the **Trevi** data set at three different stages during incremental reconstruction. Left: the initial two-frame reconstruction. Middle: an intermediate stage, after fifteen images have been added. Right: the final reconstruction with 360 photos.

Dust3R



Neural radiance field (NeRF)



NERF is slow (0.1 FPS)



Mast3r-slam

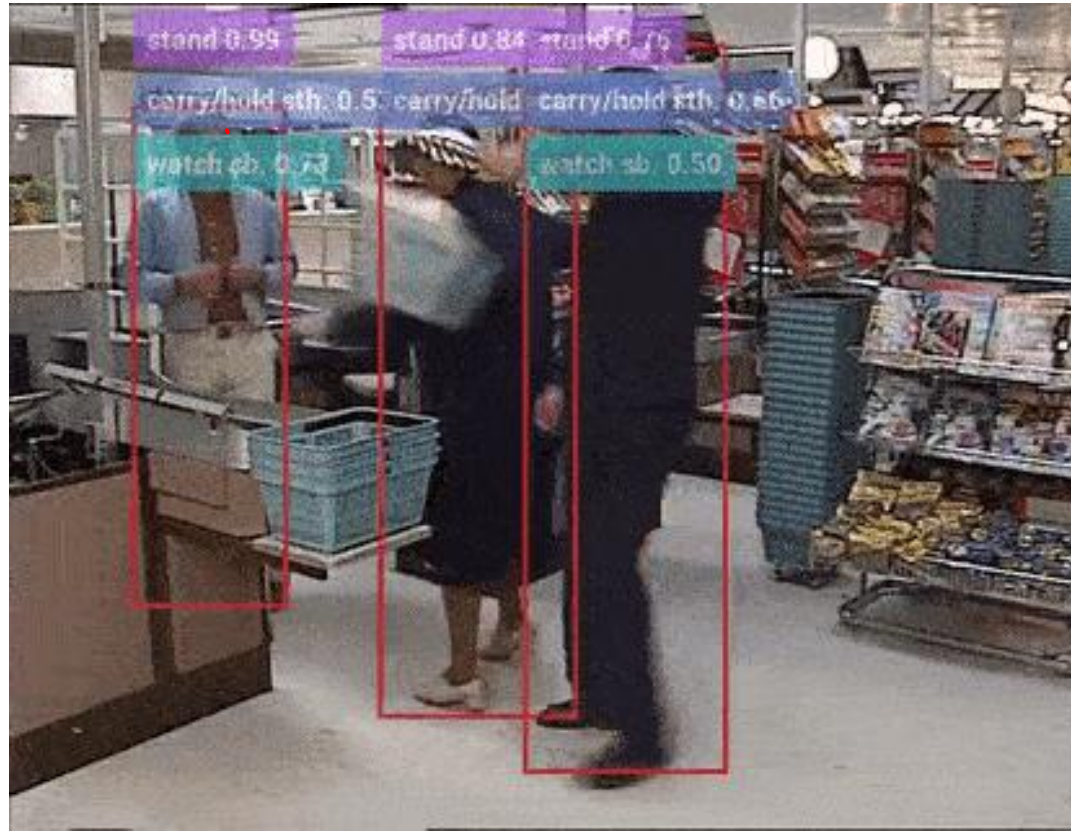



Uncalibrated RGB
8x Playback



Object tracking

1. Object detection (bounding box)
2. Object association (Hungarian association algorithm)





Hungarian algorithm

1. Row Reduction by row minima
2. Column Reduction by column minima
3. Zero Coverage
4. Iterative Refinement (if needed)
5. Finding the Optimal Assignment

The diagram illustrates the Hungarian algorithm's matrix reduction process. It features a 4x4 grid representing a cost matrix. The columns are labeled 1, 2, ..., j at the top, and the rows are labeled ID: 1, ID: 2, ..., ID: i on the left. Red circles highlight the first column and the first two rows. Red arrows indicate the subtraction of the minimum value in each row from all elements in that row. A vertical red line is drawn through the first column, and a horizontal red line is drawn through the first row, intersecting at cell C₁₁. The matrix cells are labeled as follows:

	1	2	...	j
ID: 1	C ₁₁	C ₂₁	...	C _{j1}
ID: 2	C ₁₂	C ₂₂	...	C _{j2}
...	...	...	...	...
ID: i	C _{1i}	C _{2i}	...	C _{ji}

Exercise

Solve the assignment problem online - HungarianAlgorithm.com

$1+3+4$
AD
BF

Previous **A** Previous **B** Previous **C**

$\overset{4}{1+3+0}$

Next **D**

$$1-1=\textcircled{0}$$

$$3-1=2-1=1$$

$$1-1=\textcircled{0}$$

Next **E**

$\times$ 0

$$2-1=1$$

4

Next **F**

$$6-2=4$$


$$3-2-1-1=0$$

$$2-2=0$$


Lab 3a (until 2:45 pm)

- Understand the bounding box output
- ID association between two frames
- Visualize results

Auto-encoder-decoders



x	y
1	10
2	20
5	30




$$\text{loss} = \frac{(\hat{p} - x)^2}{2}$$


$$\nabla = \frac{\partial \text{loss}}{\partial w}$$


$$w' = w - \nabla \cdot LR$$

x	y
1	10
2	20
5	30



$$\underline{CE} = -\log(p)$$






$$MSE = \frac{(y_p - y_{gt})^2}{n}$$

height
5'11"


y_{gt}
6'4"

Auto encoder decoder

loss = MSE



CNN Auto-encoder-decoder



Input


Encoder

Bottleneck

Decoder

Output

regression



mse
wte

Generating image is hard

- Multi dimensional almost infinite







Image inpainting




Old picture coloring






0.255

Style transfer



Transpose Convolution



In front of a Chinese temple (1910-1919)

1

C

B




Video generator

- Seed dance
- Google veo
- OpenAI Sora
- META Emu





x



Input

encoder
 $e_{\theta}(x)$


z



Latent Space

decoder
 $d_{\phi}(z)$

$\hat{x}$




ϵ

Output

Applications

- Feature extraction
- Data compression
- Image search
- Topic modeling
- Image coloring
- Translation
- ~~Anomaly detection~~



AI challenges

- Dataset
 - Supervised
 - Unsupervised
 - Self-supervised
- Model design
 - Data structure
 - Loss function
 - Neural network architecture

$$N \sim ((x) = y)$$


$$\text{Autoencode}(x) = x$$

Lab 3b (until 4:00 pm)

- Reconstruction
- Denoising
- Anomaly detection


Reconstruction

- Activation function
- Loss function
- Metric
- Optimizer
- Data set format




Denoising

- Activation function
- Loss function MSE
- Metric
- Optimizer
- Data set format



What if we keep adding noise




Anomaly detection

- Activation function
- Loss function
- Metric $\leftarrow \begin{matrix} MAE \\ IP \\ n \end{matrix}$
- Optimizer
- Data set format
 - Only train on normal data
 - Abnormal labels



$\left\{ \begin{matrix} 1 \\ 0 \end{matrix} \right.$



Θ




$E > \text{threshold}$ $\rightarrow \text{add mean}$

Receiver operating characteristic (ROC) area under the curve (AUC)


$$TPR = \frac{TP}{TP + FN}$$

$$FPR = \frac{FP}{FP + TN}$$




Binary classification

- High cost of false positives: low FPR (Point A), accepting a reduced TPR to minimize false alarms.
- High cost of false negatives: Choose an aggressive threshold with high TPR (Point C), accepting more false positives to avoid missed detections.
- Similar costs: Choose a balanced threshold (Point B) that trades off TPR and FPR more evenly.




Object detection metric




 **Ground Truth**

 **Prediction**



True Positive
(**TP**): $\text{IoU} \geq 50\%$

No True Negative
(**TN**)



False Negative
(**FN**): $\text{IoU} = 0$


Prediction\Ground Truth	Positive	Negative
Positive	TP	FP
Negative	FN	TN

$$accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$


$$F1 = \frac{precision * recall}{precision + recall}$$



mAP

$$precision = \frac{TP}{TP + FP} = \frac{TP}{Predictions}$$

$$recall = \frac{TP}{TP + FN} = \frac{TP}{Ground\ Truth}$$




Lab 3c (until 4:40 pm)

- Anomaly detection
- No labels for threshold

Generative Adversarial Networks (GANs)

Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... & Bengio, Y. (2020). Generative adversarial networks. *Communications of the ACM*, 63(11), 139-144.





$$\min_G \max_D (D, G) = \log(D(x)) + \log(1 - D(G(z)))$$

Generative Adversarial Networks – Creating Images

GANs can generate essentially any form of data.

For the MNIST problem, GANs can produce a flat array of 0s and 1s.



Progression of generating MNIST examples as a GAN is trained

Generative Adversarial Networks – Deep Convolution GANs


Recall that we have seen the transpose convolution. We can create a new type of generator, one that includes these transpose convolution.

DCGAN (**D**eep **C**onvolution **GAN**)

Input a random column vector

Passed through several transpose convolutions

Final result is a generated image!




Original DCGAN architecture

DCGAN – Exploring the latent space of images

With a trained DCGAN, we can manipulate the latent vector, z , to explore the space of images we can generate.

Latent Vector Space (adding the z vectors)



Pixel Space (adding the pixel values)



GAN APPLICATIONS



Figure 5: 1024×1024 images generated using the CELEBA-HQ dataset. See Appendix F for a larger set of results, and the accompanying video for latent space interpolations.




Figure 2: From left to right: bicubic interpolation, deep residual network optimized for MSE, deep residual generative adversarial network optimized for a loss more sensitive to human perception, original HR image. Corresponding PSNR and SSIM are shown in brackets. [4× upscaling]



[pix2pixHD](#)




GAN collapse






<https://link.springer.com/article/10.1007/s12065-021-00665-z>

<https://arxiv.org/pdf/1611.02163>

Insatiability



Conditional GAN (cGAN): Targeted Image generation



<https://arxiv.org/abs/1611.07004>, <https://arxiv.org/abs/1610.09585>

CycleGAN

Monet $\leftrightarrow$ Photos



Monet $\rightarrow$ photo

Zebras $\leftrightarrow$ Horses



zebra $\rightarrow$ horse

Summer $\leftrightarrow$ Winter



summer $\rightarrow$ winter




photo $\rightarrow$ Monet



horse $\rightarrow$ zebra



winter $\rightarrow$ summer



Photograph



Monet



Van Gogh



Cezanne



Ukiyo-e

CycleGAN




Figure 3: (a) Our model contains two mapping functions $G : X \rightarrow Y$ and $F : Y \rightarrow X$, and associated adversarial discriminators D_Y and D_X . D_Y encourages G to translate X into outputs indistinguishable from domain Y , and vice versa for D_X and F . To further regularize the mappings, we introduce two *cycle consistency losses* that capture the intuition that if we translate from one domain to the other and back again we should arrive at where we started: (b) forward cycle-consistency loss: $x \rightarrow G(x) \rightarrow F(G(x)) \approx x$, and (c) backward cycle-consistency loss: $y \rightarrow F(y) \rightarrow G(F(y)) \approx y$




Lab 3d

- Anomaly detection
- Discriminator
- No ground truth label?


TRAINING THE GAN

STEP 1




TRAINING THE GAN

STEP 2




TRAINING THE GAN

STEP 3




TRAINING THE GAN

BACK TO STEP 1




ANOMALY DETECTION

Input Test Point	
duration	0
src_bytes	215
service	http (1)
flag	SF (1)
src_bytes	215
.	.
.	.
.	.
dst_host_sr v_error_rat e	0



ANOMALY DETECTION

Input Test point	
duration	99999
src_bytes	3239862
service	5454
flag	SF (1)
src_bytes	215
.	.
.	.
.	.
dst_host_src_v_error_rate	0



Thank You

Q&A